

DataONE Architecture

DataONE Use Cases and APIs

Change Log

- 20090927:**
 - Split usecases into separate document
 - Added page breaks for rst2pdf
- 20090925:**
 - Initial collation of all use cases into single document.
 - Recast all diagrams to interaction diagrams using PlantUML automated tools
- 20090924:**
 - Layout document structure, draft templates

Table of Contents

Use Cases Overview	1
Use Case Implementation Priority	2
Use Cases	2
DataONE Application Programming Interfaces	41

Use Cases Overview

Date: 20090924

An important early part in the process of system architecture is identifying the functional requirements that the resulting system should be able to address. Use cases are helpful for this as they describe how actors (typically entities external to the system) interact with system components. Interaction diagrams help define the interfaces and messages transmitted between components.

Note

TODO: Architecture overview diagram TODO: Component diagram TODO: Broad categories of use cases

Use Case Layout

Each use case document is structured following a common template which includes the following elements.

Author

Who generated the use case. This will usually be "VDC TWG".

Date

Date the use case was identified and known subsequent modifications.

Version

One of "Draft", "Review", "Published"

Goal

Brief description of what is intended with the use case.

Summary

An overview of the use case which includes the goal and major actors.

Use case diagram

A diagram representation of the use case, somewhat following the UML.

Sequence diagram

A sequence diagram that shows the operations involved in addressing the use case.

Actors

The participants in the use case - includes people and systems.

Preconditions

The conditions that must exist for the use case. For example, a particular use case may only be valid if the use has been authenticated.

Triggers

Describes the even that causes the use case to be initiated.

Postconditions

The state of the system after the use case - what conditions are true after the use case is completed.

Notes

Miscellaneous notes and discussion items that can't be placed in the various categories. Comments should be annotated with time and author.

Use Case Implementation Priority

Prioritizing system implementation to address the use cases involves several factors including:

- *Vision of the project.* The system is being designed with some overall goals described by the vision of the project proposal.
- *Requirements of the community.* The stakeholders that comprise the user and participant community quite likely has some opinion on functionalities of the system that are important to them. If these are not properly addressed, then the resulting system may appear as a failure to them.
- *Requirements of the sponsor.* The sponsor has laid out goals in the RFP that the project is responding to and also in the final agreement for conducting the work.
- *Dependencies between use cases.* Implementation of functionality to address some use cases will require implementation of some components not directly specific in a use case.
- *Resources available for implementation.* Some use cases may be identified as high priority, but would require resources that would prevent implementation of a number of other lower priority features.

Use Cases

Use Case 1 - CRUD Read()	4
Use Case 02 - List GUIDs By Search	6
Use Case 03 - Register MN	8
Use Case 04 - CRUD (Create, Update Delete)	10
Use Case 05 - CRUD (Create, Update Delete)	11
Use Case 06 - MN Synchronize	12
Use Case 07 - CN Batch Upload	13
Use Case 08 - Replication Policy Communication	14
Use Case 09 - Replicate MN to MN	15
Use Case 10 - MN Status Reports	16
Use Case 11 - CRUD Workflow Objects	17
Use Case 12 - User Authentication	18
Use Case 13 - User Authorization	19

Use Case 14 - System Authentication and Authorization	20
Use Case 15 - Account Management	21
Use Case 16 - Log CRUD Operations	22
Use Case 17 - CRUD Logs Aggregated at CNs	23
Use Case 18 - MN Retrieve Aggregated Logs	24
Use Case 19 - Retrieve Object Download Summary	25
Use Case 20 - Owner Retrieve Aggregate Logs	26
Use Case 21 - Owner Subscribe to CRUD Operations	27
Use Case 22 - Link/Citation Report for Owner	28
Use Case 23 - Owner Expunge Content	29
Use Case 24 - MNs and CNs Support Transactions	30
Use Case 25 - Detect Damaged Content	31
Use Case 26 - Data Quality Checks	32
Use Case 27 - Metadata Version Migration	33
Use Case 28 - Derived Product Original Change Notification	34
Use Case 29 - CN Load Balancing	35
Use Case 30 - MN Outage Notification	36
Use Case 31 - Manage Access Policies	37
Use Case 32 - Transfer Object Ownership	38
Use Case 33 - Search for Data	39
Use Case 34 - CNs Support Other Discovery Mechanisms (e.g. Google)	40

Use Case 1 - CRUD Read()

Author

VDC TWG

Date

- 20090924 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft

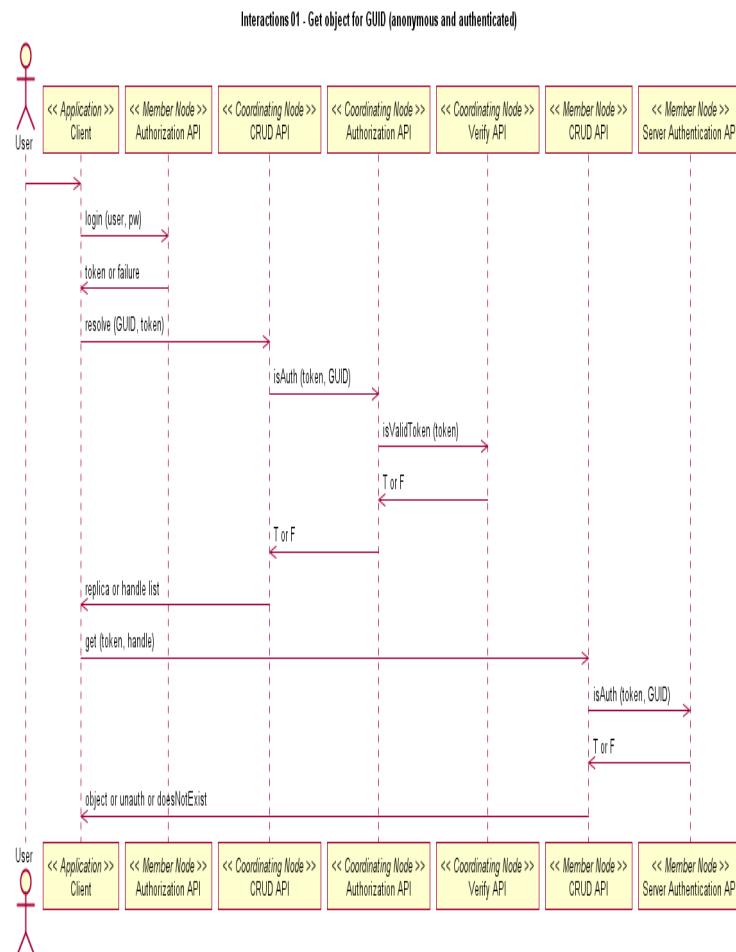
Goal

Get object identified by GUID (authenticated or not, notify subscriber of access).

Summary

A client has an identifier for some object within the DataONE system and is retrieving the referenced object. The DataONE system must resolve the identifier and return the object bytes after checking that the user has read privileges on the object.

Watchers (subscribers) of the object are notified of the retrieval (what are they notified of?).



Actors

- Client requesting object
- Coordinating Node
- Member Node

Preconditions

- Client has authenticated to at the desired level (e.g. client may not have authenticated, so access might be anonymous).

Triggers

- Watchers are notified of the retrieval
- Operation is logged (are loggers a type of watcher?)

Post Conditions

- The client has a copy of the object bytes (or an error message in the case of failure)
- The log is updated with the results of the operation
- Watchers perform their processing

Notes

- for the GET operation, the coordinating node does not need to check access, because the member node will do this check before delivering the results.
- Data sent to watchers might include: timestamp, object identifier, user id, IP of client,

Use Case 02 - List GUIDs By Search

Author

VDC TWG

Date

- 20090927 Provided more detail to use case
- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications

Version

Draft

Goal

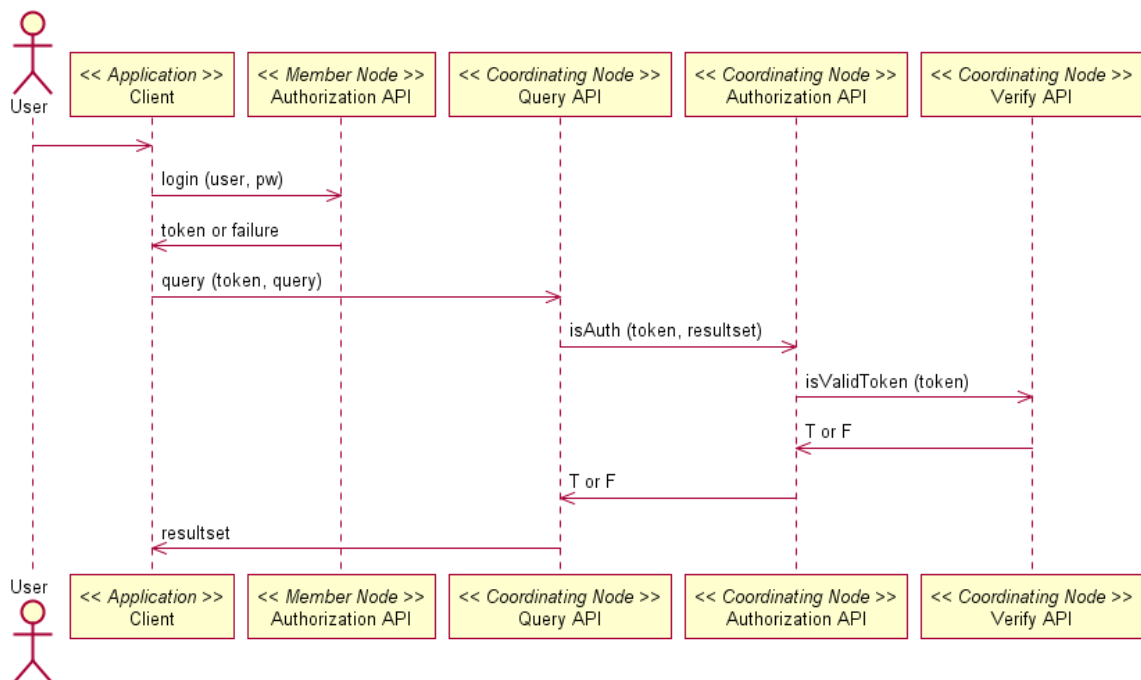
Get list of GUIDs from metadata search (anon and authenticated).

Summary

A user performs a search against the DataONE system and receives a list of object identifiers (GUIDs) that match the search criteria. The list of GUIDs is filtered such that only objects for which the user has read permission will be returned.

This use case assumes that the search is being performed by submitting a query against a CN.

Interactions 02 - Client query to get list of GUIDs from metadata search



Actors

- Client performing search operation
- Coordinating Node

Preconditions

- Client has authenticated to at the desired level (e.g. client may not have authenticated, so access might be anonymous).

Triggers

- Operation is logged

Post Conditions

- The client has a list of GUIDs for which they have permission to read and match the supplied query, or an error message if the query was invalid.
- The log is updated with the request

Notes

- Two #2's on diagram. One to Authentication API and one to query API.
- lots of stuff buried in query. What information/interface is needed for resolution/expansion of the query? For example a particular user may be identified with a particular community which uses a particular ontology. Where are things like ontologies maintained and where are the interfaces between these and searches maintained? Searches need to be able to respond to the dynamic nature of domain ontologies. Does the ontology need to have authorization tied to it? A use case would be an ontology which is under development and not ready for general distribution. Another case might be where the ontology contains representations of human subject information, meaning that "cancer survivors" includes this list of people. Or where the ontology provides geospatial information about restricted information (locations of USFS FIA plots). Similar issues apply to T&E species information.
- Group structure representation: A user is a member of zero to many groups, authorization may be carried with that group membership. This is down in the verify API, which needs to take into account group membership and group permission inheritance. Question: What is the resolution for conflicting access permissions?
- Can a query be done at a member node? The diagram implies that the query is only done at a coordinating node. A member node may have a subset (or the full set) of metadata and the user may wish to do the search only at a member node, knowing that the subset is restricted to data relevant to a community of practice.
- Where is the filtering done for the results, in terms of when does the result set get filtered to those which the user can see. Does the model allow the query service to do two searches and union the results? The first query might be against the set of records known to be publicly readable, with a follow-on query against the set of records known to have restricted access. Is this envisioning the query service generating a set of results (against the restricted records) and passing that list to the authorization API, or is this envisioning a one by one situation. Or does the authorization API pass back to the query service the list of records that the user is authorized to see, and that gets anded with the result set.
- How distributed is the database of users and their authorization records? Does the coordinating node have to know all authorization records, with that being replicated? Or is some amount of the authorization maintained by (for example) the "owning" member node and the search system may need to contact multiple authorization authorities for the records in a particular result set? It also implies that if the authorization information is at the coordinating node, that the authorization authority needs a means to validate the replicated authorization information to ensure that the info at the cn is up-to-date.
- Why does step 4 need to validate the token, given that the token was generated in step 2? Answer is that this is normal security validation of all user inputs once those have gone out of the control of the issuing service. Also applies to cases where tokens have a specific lifetime, and the token may be valid, but expired. Authorization token lifetimes are important for both interactive sessions and for cases like workflows, where the session may run for days.
- What logging level is required for the operation? Should the search appear as a hit against each object? Just log the query (and perhaps who did it)?

Use Case 03 - Register MN

Author

VDC TWG

Date

- 20090927 More detail added to use case description
- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft

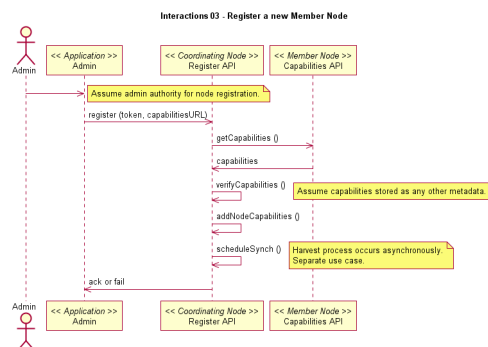
Goal

Register a new Member Node.

Summary

This use case describes the technical process for addition of a new member node (MN) to the DataONE infrastructure. It is assumed that the appropriate social contracts have been formed and the MN is operational, ready to be connected.

The MN is identified by a URL which is the service endpoint. A DataONE administrator adds the MN URL to the MN registry. A CN retrieves the registration request from the message queue and queries the MN for capabilities, then verifies those capabilities match what was advertised. If everything is OK, the MN is made a live member of DataONE and replication of content on the MN is scheduled. The new MN also becomes a receiver for content replicated from other MNs.



Actors

MN, CN, Administrator

Preconditions

- The MN is operational
- There is an agreement between the DataONE operators and the MN operator that the MN is to be added to DataONE
- The CNs are ready for receiving MN registrations

Triggers

- Notification is sent to MN administrators of the operation status
- Content replication from the new MN is initiated
- The new MN is added to the list of targets for replication
- The system logs are updated

Post Conditions

- The new MN operates as part of the DataONE infrastructure
- DataONE infrastructure resources are incremented by the amount available at the new MN
- The operation outcomes are logged

Notes

- Specify default replication policies
- Also check for version updates

- Should new nodes be registered with specified trust levels?
- Data providers that use acceptable services can still be discovered and accessed without registering in the DataONE registry (Conflicting, as someone has to register it)
- Allow service providers to register their services, (such as data extraction services) (ala GEOSS), but include mapping to higher semantic model
- For well known services, registration system must be able to describe constraints (e.g., allowable inputs, outputs, algorithms that can be specified)
- Services can be parameterized
- Member node should be able to request the coordinating node to re-validate its capabilities list.
- Are there different levels of trust for member nodes?

Use Case 04 - CRUD (Create, Update Delete)

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

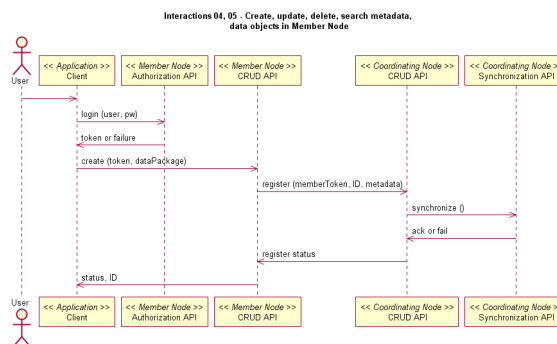
Draft

Goal

Create, update or delete metadata record on a Member Node.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Are IDs are pegged to versions? If so, update might differ more radically from create.
- If delete is a metadata operation only, then will be very similar to update. But if it is a physical operation, it becomes more complex. Will there be an archival copy? Will member nodes that hold replicas be "forced" to destroy the replicas?
- Is submit from a member node a push or pull?
- Should queuing of a submission within a member node be required for success? (We assume yes, to guarantee replication of submissions).

Use Case 05 - CRUD (Create, Update Delete)

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

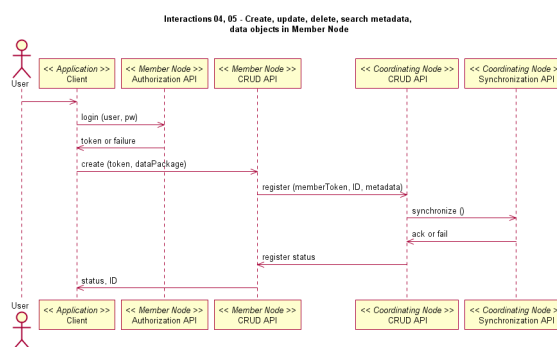
Draft

Goal

Create/Update/Delete data object in Member Node.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- See all questions above (Use Case #4) for metadata operations.
- client should be able ask CN for a list of acceptable MNs to upload data with specific characteristics (e.g., size, institution, theme)

Use Case 06 - MN Synchronize

Author
VDC TWG

Date


- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version
Draft

Goal
Replicate / synchronize metadata record between Member Node and Coordinating Node.

Summary

XXX



Actors
XXX

Preconditions
XXX

Triggers
XXX

Post Conditions
XXX

Notes

XXX

Use Case 07 - CN Batch Upload

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

Batch Operations - Coordinating Node requests metadata /data list from new Member Node and then batch upload (disable indexing for example to improve insert performance).

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 08 - Replication Policy Communication

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

Communication of replication policy metadata among Member Nodes and Coordinating Nodes.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 09 - Replicate MN to MN

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

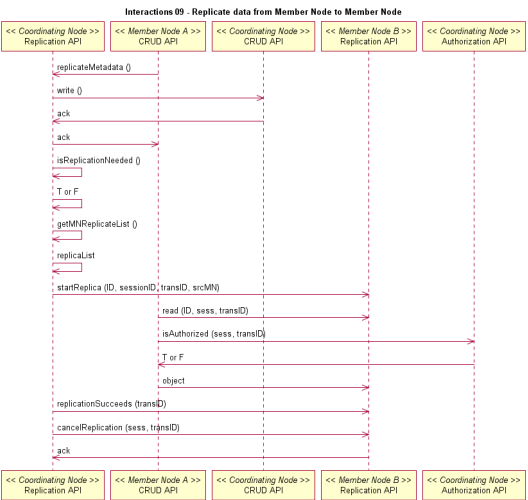
Draft

Goal

Replicate data from Member Node to Member Node - (facilitated by Coordinating Node or independently but notifying Coordinating Node of operation).

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 10 - MN Status Reports

Author
VDC TWG

Date

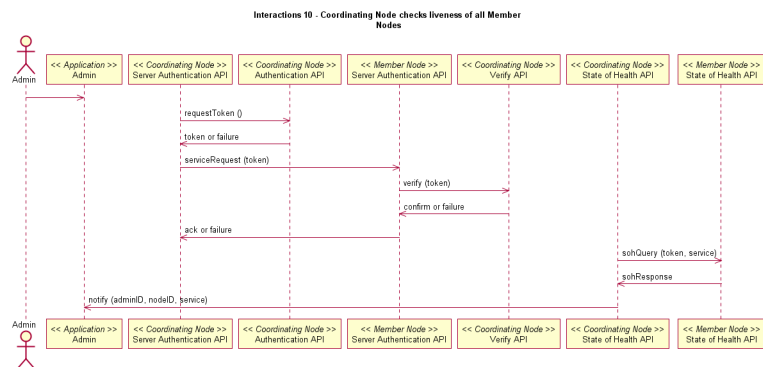
- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version
Draft

Goal
Coordinating Node checks "liveness" of all Member Nodes - checks ping, service x, load, space, bandwidth, transaction rate, ...

Summary

XXX



Actors
XXX

Preconditions
XXX

Triggers
XXX

Post Conditions
XXX

Notes

XXX

Use Case 11 - CRUD Workflow Objects

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

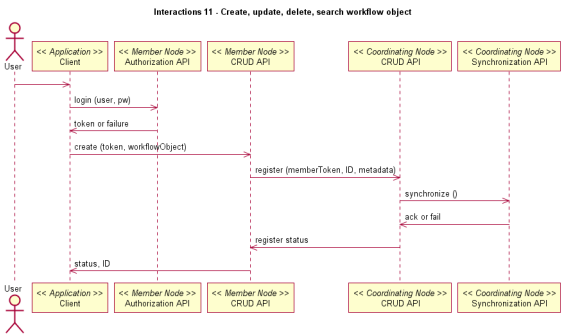
Draft

Goal

Create / update / delete / search workflow objects.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 12 - User Authentication

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft


Goal

User Authentication - Person via client software authenticates against Identify Provider to establish session token.

Summary

XXX

Interactions 12 - Person authenticates against Identity Provider



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- To be a bit more abstract, it might be good to use something like login(user, credentials) or login(subject) for the method signature. Here "subject" would be similar to the [javax.security.auth.Subject](#) concept. (PEA)
- Can logins also happening against a MN, which simply proxies the request up to the CN? (PEA)

Use Case 13 - User Authorization

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

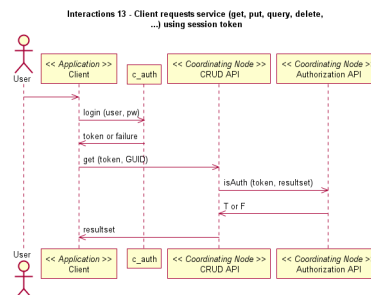
Draft

Goal

User Authorization - Client requests service (get, put, query, delete, ...) using session token.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- It would be great to factor out the authentication steps (#1, 2) from the diagram and replace it with some sort of precondition like "Client has authentication token." (PEA)

Use Case 14 - System Authentication and Authorization

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

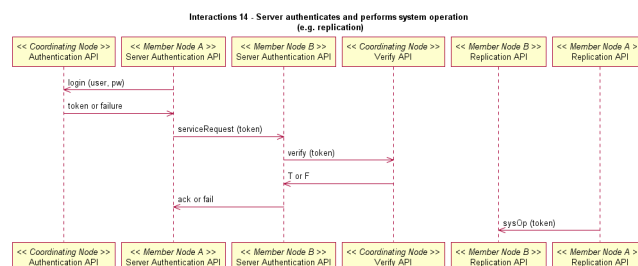
Draft

Goal

System Authentication/Authorization - Server authenticates and performs system operations (e.g. replication).

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- I don't understand Step #7. To me the serviceRequest in step #3 is the actual request for MN-B to do a "system operation", which kicks off the verification steps #4 and #5. If this example service request is replication, then step #3 should point to the replication API, and internally MN-B decides that it must perform the verify step #4. Then step #6 is the response to the system operation, and step #7 is not needed. (I think that this diagram must be showing a kind of pre-approval process for the operation in step #7, but it doesn't make sense to me because in step #7, MN-B will just have to re-verify the token again.) (PEA)

Use Case 15 - Account Management

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

User Account Management - Create new user account on Identity Provider (also edit, delete, ...).

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- This diagram should include a step for email confirmation.
- By default, accounts have no real privileges. To get higher privileges, users may have to jump through more hoops (such as verifying their association with a project/institution)
- Presumably, if we are using external identity providers this user account management functionality isn't provided by the CN. Right? (PEA)

Use Case 16 - Log CRUD Operations

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft


Goal

All CRUD operations on metadata and data are logged at each node.

Summary

XXX

Interactions 16 - All CRUD operations on metadata and data are logged at each node



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 17 - CRUD Logs Aggregated at CNs

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

All CRUD logs are aggregated at coordinating nodes.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 18 - MN Retrieve Aggregated Logs

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

Member nodes can request aggregated CRUD log for {time period/object id/userid} for all of 'their' objects.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 19 - Retrieve Object Download Summary

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

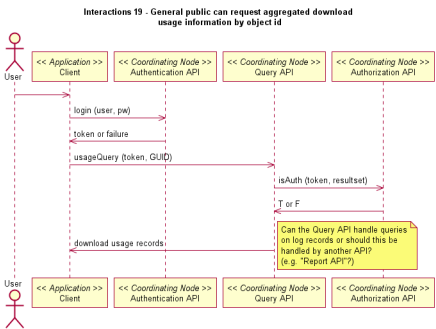
Draft

Goal

General public can request aggregated download usage information by object id.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 20 - Owner Retrieve Aggregate Logs

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

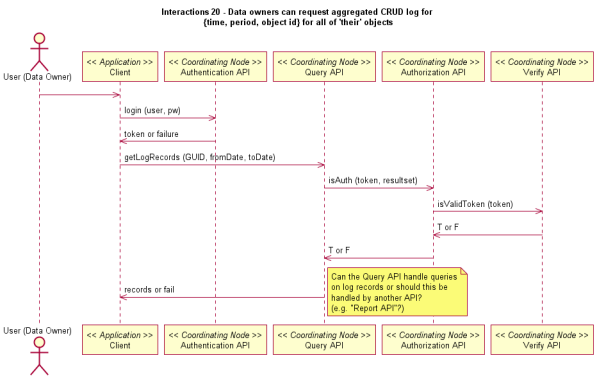
Draft

Goal

Data owners can request aggregated CRUD log for {time period/object id} for all of 'their' objects.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 21 - Owner Subscribe to CRUD Operations

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

Data owners can subscribe to notification service for CRUD operations for objects they own.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 22 - Link/Citation Report for Owner

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

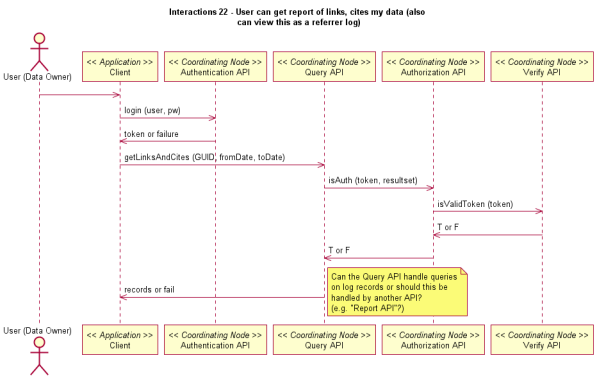
Draft

Goal

User can get report of links/cites my data (also can view this as a referrer log).

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 23 - Owner Expunge Content

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

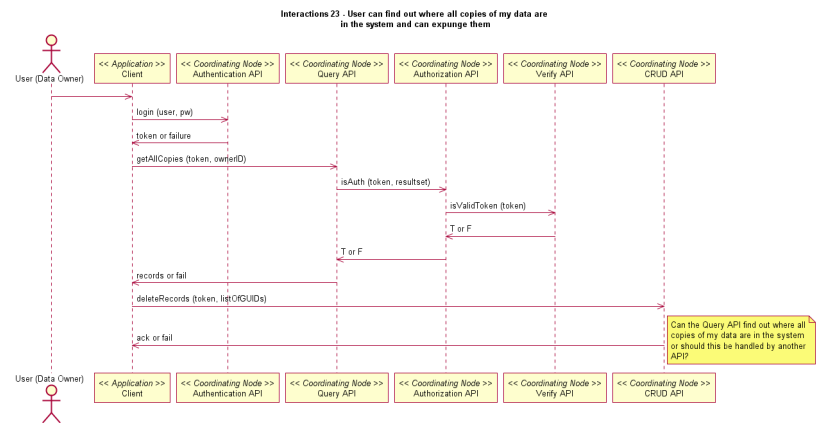
Draft

Goal

User can find out where all copies of my data are in the system and can expunge them.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Dealing with inappropriate usage of (ie, copyrighted objects that are stored)

Use Case 24 - MNs and CNs Support Transactions

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft

Goal

Transactions - CNs and MNs should support transaction sets where operations all complete successfully or get rolled back (e.g., upload both data and metadata records).

Summary

XXX

Interactions 24 - CNs and MNs should support transaction sets

Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Does this include just one service, or sets of services that can be included in a transaction?
- Can transactions span multiple MNs? MNs and CNs?

Use Case 25 - Detect Damaged Content

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

System should scans for damaged/defaced data and metadata using some validation process.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 26 - Data Quality Checks

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

System performs data quality checks on data.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Is this more for verifying the quality of metadata? Otherwise how does it differ from #25?

Use Case 27 - Metadata Version Migration

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft

Goal

CN should support forward migration of metadata documents from one version to another within a standard and to other standards.

Summary

XXX

Interactions 27 - Migration -- CN should support forward migration of metadata documents from one version to another within a standard and to other standards

Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Should this apply to data as well?

Use Case 28 - Derived Product Original Change Notification

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

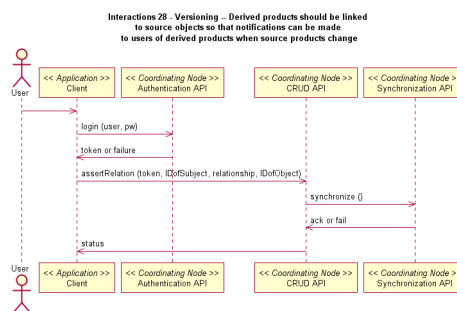
Draft

Goal

Relationships/Versioning - Derived products should be linked to source objects so that notifications can be made to users of derived products when source products change.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Who can assert these relationships? The relevant object owners? Anybody?
- How are relationships removed?

Use Case 29 - CN Load Balancing

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft

Goal

Load Balancing - Requests to coordinating nodes are load balanced.

Summary

XXX

Interactions 29 - Load balancing between Coordinating Nodes

Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 30 - MN Outage Notification

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft


Goal

MN can notify CN about pending outages, severity, and duration, and CNs may want to act on that knowledge to maintain seamless operation.

Summary

XXX

Interactions 30 - MN can notify CN about pending outages, severity and duration, and CNs may want to act on that knowledge to maintain seamless operation



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 31 - Manage Access Policies

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

Manage Access Policies - Client can specify access restrictions for their data and metadata objects. Also supports release time embargoes.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Users can be members of groups that can participate in access directives.
- I have removed the phrase "and replication" from the use case statement because use case #8 deals with setting replication policies. (PEA)
- Step #1, should have a signature of setAccess(token, GUID, accessPolicy). Even though the diagram says "Assume user authority for specifying restrictions", practically speaking we will need to verify that authority and the user's identify with a token. Also "accessLevel" sounds very limited, and access policy implies a possibly more sophisticated access policy delineation, including embargoes.

Use Case 32 - Transfer Object Ownership

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

User or organization takes over 'ownership' of a set of objects (write access for orphaned records)

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

- Is there an authority for any given data/metadata object?

Use Case 33 - Search for Data

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version


Draft

Goal

Clients should be able to search for data using CN metadata catalogs.

Summary

XXX



Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

Use Case 34 - CNs Support Other Discovery Mechanisms (e.g. Google)

Author

VDC TWG

Date

- 20090925 Adding to use case description and documentation
- 20090114 (TWG meeting). Subsequent various modifications.

Version

Draft

Goal

Coordinating Nodes publish metadata in formats for other discovery services like Google/Libraries/GCMD/etc.

Summary

XXX

Interactions 34 - Discovery; Coordinating Nodes publish metadata in formats for other discovery services like Google, Libraries, GCMD, etc.

Actors

XXX

Preconditions

XXX

Triggers

XXX

Post Conditions

XXX

Notes

XXX

DataONE Application Programming Interfaces

isAuthorized (sess, transID) read (ID, sess, transID)

Member Node APIs

Authorization API

```
login (user, pw)
isAuth (token, GUID)
```

CRUD API

```
batchGet (objectsList)
checkDataQuality (GUID)
create (token, dataPackage)
create (token, workflowObject)
get (token, handle)
getChecksum (GUID)
getLogRecords (fromDate, toDate)
log (create, creationDate, ID)
register (memberToken, ID, metadata)
```

Replication API

```
cancelReplication (sess, transID)
listAllObjects ()
listRecords ()
replicationSucceeds (transID)
startReplica (ID, sessionID, transID, srcMN)
sysOp (token)
```

Server Authentication API

Capabilities API

```
serviceRequest (token)
verify (token)
getCapabilities ()
State of Health API
sohQuery (token, service)
```

Coordinating Node APIs

Authentication API

```
login (user, pw)
requestToken ()
newAccount (user, pw)
```

Verify API

```
isValidToken (token)
Authorization API
isAuth (token, GUID)
isAuth (token, resultset)
isValidToken (token)
setAccess (GUID, accessLevel)
takeOwnership (objectList)
```

Query API

```
getAllCopies (token, ownerID)
getLinksAndCites (GUID, fromDate, toDate)
getLogRecords (GUID, fromDate, toDate)
getLogRecords (nodeID, fromDate, toDate)
query (token, query)
searchForData ()
usageQuery (token, GUID)
```

CRUD API

```
assertRelation (token, IDofSubject, relationship, IDofObject)
batchWrite (objectsList)
deleteRecords (token, listOfGUIDs)
get (token, GUID)
register (memberToken, ID, metadata)
resolve (GUID, token)
storeRecords ()
write ()
```

Replication API

```
getDefaultReplicationPolicy (sess)
getMNReplicateList ()
getReplicationPolicy (sess, ID)
isReplicationNeeded ()
replicateMetadata ()
setDefaultReplicationPolicy (sess, policy)
setReplicationPolicy (sess, policy, ID)
```

Synchronization API

```
pollForNewIDs ()  
synchronize ()  
verifyAllRecordsReceived ()
```

Register API

```
addNodeCapabilities ()  
register (token, capabilitiesURL)  
registerForNotification (token, dataOwnerID)  
scheduleSynch ()  
verifyCapabilities ()
```

Server Authentication API

State of Health API

```
actOnPendingEvent (event)  
reportPendingEvent (event)  
Admin Application  
notify (adminID, nodeID, service)
```